

David Isaac Belais

505 SE 35th Ave, Portland OR 97214 | 503-267-0942 | david@belais.me | david.belais.me | github.com/davebelais

Summary

I am a highly productive data engineer and software engineer with 18 years of relevant experience.

I pride myself in:

- Creating resilient, maintainable, integrous data pipelines and products (ET/ELT batch, micro-batch, and streaming for both OLAP and OLTP databases/lake houses/data lakes/data warehouses) for "big data", master data, fact data and dimensional data
- Authoring elegant, bulletproof, type-annotated, well-formed, thoroughly tested, distributable Python libraries, CLIs, asynchronous micro-services (web APIs), and SDKs
- Designing efficient, maintainable, testable, continuously integrated and deployed, modern software systems (CI/CD, Test-driven development)
- Planning development work with clarity, flexibility, parallel execution, and collaboration in mind (whether using Agile or waterfall)
- Leading engineering teams with complex and ambiguous directives towards clear, executable road maps
- Condensing fact from the vapor of nuance while maintaining traceability and continuity with data provenance, and facilitating data governance
- Utilizing "AI" and machine learning tools and frameworks *selectively, prudently*, and with traceability and long-term *cost efficacy* in mind

Skills

I have professional experience with (not exhaustive):

- Platforms: Databricks, Snowflake, Amazon Web Services
- Languages: Python, SQL, C++, Javascript, HTML, XML, PHP, WSDL, Rust (AWS - including Lambda, EMR, Aurora, IAM, Cloudformation, EC2, S3)
- Databases and query engines: Databricks Lakehouse, Deltalake, Snowflake, Terradata, Netezza, Hive, Presto, DuckDB, PostgreSQL, MySQL, SQL Server, Oracle, IBM DB2, SQLite
- Applications, Services and Frameworks: Apache Spark, Apache Kafka, SQLAlchemy, FastAPI, Flask, Docker, Terraform, Linux, Unix, Github Actions, Jenkins, Hadoop, Copilot
- Protocols and Specifications: Open API (Swagger), SOAP, MIME, AS2 (for GDSN data pools), ASGI, WSGI
- Distributed File Systems: DBFS, S3, HDFS

Experience

Nike | Lead Data Engineer - Sustainability Analytics | March 2021 - June 2025

Platforms: Databricks, Snowflake, Amazon Web Services (AWS - including Lambda, EMR, Aurora, IAM, Cloudformation, EC2, S3)

Languages: Python, SQL, Javascript

Databases and query engines: Databricks Lakehouse, Deltalake, Snowflake, Terradata, Hive, Presto, PostgreSQL, Oracle, SQLite

Applications, Services and Frameworks: Apache Spark, Apache Kafka, SQLAlchemy, Alembic, FastAPI, Terraform, Docker, Linux, Github Actions, Jenkins, Hadoop, Copilot

Protocols and Specifications: Open API (Swagger), ASGI

Distributed File Systems: DBFS, S3, HDFS

- I lead a team of, variably, 4-8 data engineers in developing data and software products supporting analysts, data scientists, environmental scientists, product developers, and sustainability professionals in assessing and mitigating Nike's environmental impacts
- I redesigned and implemented data pipelines leveraging Databricks Deltalake, Snowflake, and Amazon EMR—employing batch, micro-batch, and streaming (Apache Kafka and Spark streaming, both as producer and consumer) and Delta live tables, reducing aggregate compute costs by 80% as compared with legacy pipelines
- I authored asynchronous Python web APIs using FastAPI and SQLAlchemy on AWS Lambda, using Terraform for infrastructure as code, and OAuth2 (Okta) for authentication, to facilitate preemptive mitigation of environmental impacts by facilitating pre-manufacture scenario modeling in product development systems
- I designed our systems for calculating material and product footprints as individually testable component libraries, permitting us to fully employ test-driven development, and thereby safely make use of continuous integration and deployment (CI/CD), permitting us to frequently release multiple features daily
- I leveraged data models in our Databricks Deltalake, Snowflake databases, and (prior to 2023) S3/hive data lake incorporating type 2 slowly changing dimensions in order to address obstacles to replicating historically reported metrics (required for regulatory audits)
- I authored foundational data products exposing Environmental Health & Safety data from our 3rd-party EHS reporting system, Enablon (ingested via their "Blink" OData API), into our Databricks Deltalake, Snowflake, and (prior to 2023) S3 + hive/presto data lake
- I authored enterprise developer tools including CLIs (command line interfaces) and frameworks (Python libraries) for use in CI/CD and locally, numerous SDKs (software development kits) for internal and third-party platforms, extended SQLAlchemy and Alembic to facilitate use of ORMs (object relational mappings) across multiple dialects simultaneously, and to facilitate common and complex data frame operations in Spark, validate data products based on ORM metadata, securely retrieve managed credentials, and many other common development tasks

BICP @ Nike | Lead Data Engineer - Sustainability Analytics | March 2020 - March 2021

Platforms: Snowflake, Amazon Web Services (AWS - including Lambda, EMR, Aurora, IAM, Cloudformation, EC2, S3)

Languages: Python, SQL, Rust

Databases and query engines: Snowflake, Terradata, Hive, Presto, PostgreSQL, Oracle, SQLite

Applications, Services and Frameworks: Apache Spark, SQLAlchemy, Alembic, Docker, Linux, Terraform, Github Actions, Jenkins, Hadoop

Protocols and Specifications: Open API (Swagger)

Distributed File Systems: S3, HDFS

I lead a team of data engineers in building foundational data products supporting sustainability initiatives:

- I developed a SQLAlchemy-ORM-based framework for automating deployment and versioning (schema migration) supporting all database dialects leveraged by the Nike Enterprise Data & Analytics organization: Databricks, Snowflake, Hive/Presto on S3, and PostgreSQL with full rollback and versioning support.
- I authored a framework for Sustainability Analytics' ETL jobs incorporating end-to-end schema-based data validations, local testing, and environment and file system abstraction.

BICP @ Nike | Senior Data Engineer - Sustainability Analytics | January 2020 - March 2020

Platforms: Snowflake, Amazon Web Services (AWS - including Lambda, EMR, Aurora, IAM, Cloudformation, EC2, S3)

Languages: Python, SQL, Rust

Databases and query engines: Snowflake, Terradata, Hive, Presto, PostgreSQL, Oracle, SQLite

Applications, Services and Frameworks: Apache Spark, SQLAlchemy, Alembic, Docker, Linux, Jenkins, Hadoop

Protocols and Specifications: Open API (Swagger), ASGI

Distributed File Systems: S3, HDFS

Infrastructure as Code: Terraform

The Kroger Co. | Lead Data Engineer - Web & Digital Analytics | May 2018 - November 2019

Languages: Python, SQL, Javascript, HTML, XML, WSDL

Databases and query engines: Netezza, Hive, Presto, SQL Server, IBM DB2, SQLite

Applications, Services and Frameworks: SQLAlchemy, Flask, Hadoop, Magento Commerce, IBM Websphere Commerce

Protocols and Specifications: Open API (Swagger), SOAP, MIME, AS2 (for GDSN data pools), WSGI

I lead development of:

- Data products distilling and exposing analytics to buyers and planners correlating digital and store sales and EBITDA with inventory,sell-through, prices, and promotional events—contributing to decisions resulting in a 56% increase in e-commerce sales in 2018 vs 2017, and a 67% increase in ecommerce sales in 2019 vs 2018.
- Pricing/promotions and product information integration services for Magento Commerce.

The Kroger Co. | Lead Data Engineer - Product Information Management | November 2013 - May 2018

Languages: Python, SQL, Javascript, HTML, XML, WSDL

Databases and query engines: Netezza, Hive, Presto, SQL Server, IBM DB2, SQLite

Applications, Services and Frameworks: SQLAlchemy, Flask, Hadoop

Protocols and Specifications: SOAP, MIME, AS2 (for GDSN data pools), WSGI

- I lead development of multi-platform (Spark/Hive/Presto, Netezza, DB2, Python, SQL Server, SQLAlchemy) OLAP and OLTP data products to ingest, consolidate and normalize sales, dimensional, and click-stream data from disparate subsidiary and partner systems' transactional databases, streaming platforms, APIs, and mainframes.
- I engineered algorithms for scoring semi-structured data and performing human-in-the loop data validation and auditing for product descriptions, specifications and photography acquired through trading partners (Python, SQL).
- I established source-management capabilities for inbound data to handle complex retailer/vendor/manufacter relationships (Python, SQL).
- Collaborated with emerging digital initiatives to ensure the capture of all metrics needed to facilitate accountability and continuous operational improvement.

The Kroger Co. (Fred Meyer Stores Inc.) | Business Systems Analyst - Ecommerce | March 2011 - November 2013

Languages: Python, SQL, Javascript, HTML, XML, WSDL

OLTP Databases: SQL Server, IBM DB2, SQLite

Applications, Services and Frameworks: SQLAlchemy

Protocols and Specifications: SOAP, MIME, AS2 (for GDSN data pools)

- I researched, designed, and prototyped Fred Meyer's (and later Kroger's) product information management system for customer-facing digital initiatives.
- I collaborated with Fred Meyer's technology partner, 1WorldSync, to establish a roadmap, data model, and procedures for sourcing and validating product data from GDSN data pools for use in digital sales channels.

Dissent Graphics Inc. | Full-Stack Developer | January 2008 - March 2011

Languages: Python, PHP, SQL, Javascript, Actionsript

- I designed and developed web applications for clients including: The Garrigan Lyman Group, Microsoft, Best Buy, Avenue A Razorfish, Nereus Communications, BlackEyedPeas.com, TeeFury.com, the Travel Channel's Man v. Food, TheWho.com, Custom Rights, Hello Minor, ExoticTravelers.com, and the ACLU of Oregon.

Education

- Portland State University | Computer Science (Postbaccalaureate) | 2018
- The Art Institute of Portland | Media Arts & Animation | Bachelor of Science | 2007
- Portland State University | Web Design | 2002 - 2004
- Loyola Marymount University | Fine Arts | 2000 - 2001

Open Source Projects

...because code examples are worth a thousand interview questions!

- [git-author-stats](#): A CLI and library for extracting periodic author "stats" (insertions and deletions) for a Git repository or Github organization
- [dependence](#): A CLI and library for aligning a python projects' declared dependencies with the package versions installed in the environment in which dependence is executed, and for "freezing" recursively resolved package dependencies (like pip freeze, but for a package, instead of the entire environment).
- [maya-zen-tools](#): An Autodesk Maya extension providing modeling tools for manipulating a polygon mesh using dynamically created NURBS curves and surfaces to distribute vertices and/or UVs
- [oapi](#): A python library for generating client SDKs from Open API documents
- [gittable](#): A CLI and library for performing common, but complex, development and CI/CD tasks for a Git repository, such as tagging a commit with your current project/package version and downloading or accessing specific file(s) from a remote repository (including non-public repos)

Please see [github.com/davebelais](#) and [github.com/enorganic](#) for additional code examples.

Certifications

- JQL for Admins
- Jira Automation for Admins
- AWS Certified Big Data - Specialty
- AWS Certified Cloud Practitioner